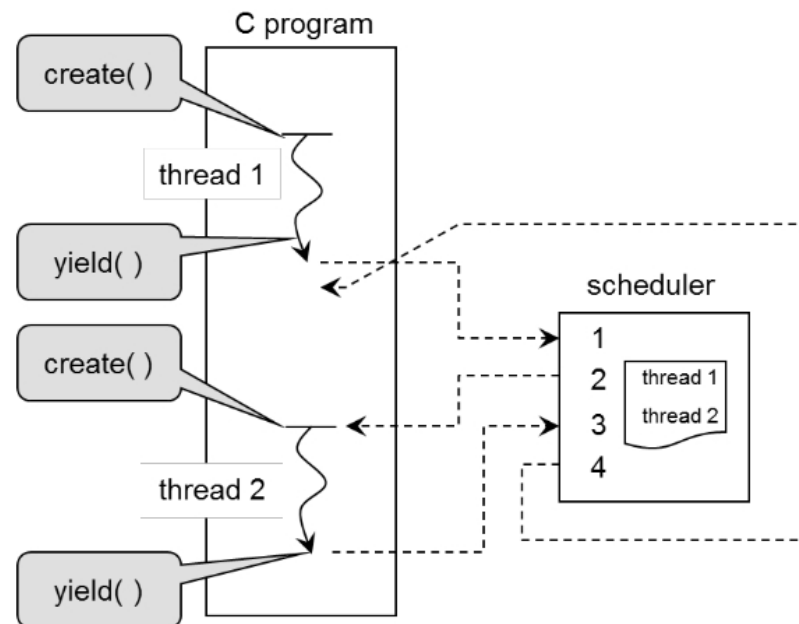


## Mapping DFGs to Single Processors: Multi-Thread Dynamic Schedule

In multi-threaded programming, each *actor* (implemented as a function) lives in a separate thread

The threads are time-interleaved by a scheduler in single processor environments

Systems in which threads **voluntarily** relinquish control back to the scheduler is referred to as *cooperative* multithreading



Such a system can be implemented using two functions *create()* and *yield()* as shown

**Mapping DFGs to Single Processors: Multi-Thread Dynamic Schedule**

The scheduler can apply different strategies to schedule threads, with the simplest one shown above as a *round-robin* schedule

*Quickthreads* is a **cooperative multithreading** library

The quickthreads API (Application Programmers Interface) consists of 4 functions

- *spt\_init()*: initializes the threading system
- *spt\_create(stp\_userf\_t \*F, void \*G)* creates a thread that will start execution with user function *F*, and will be passed a single argument *G*
- *spt\_yield()* releases control over the thread to the scheduler
- *spt\_abort()* terminates a thread (prevents it from being scheduled)

Here's an example

```
#include "../qt/stp.h"  
#include <stdio.h>
```

**Mapping DFGs to Single Processors: Multi-Thread Dynamic Schedule**

```
void hello(void *null)
{
    int n = 3;
    while (n-- > 0)
    {
        printf("hello\n");
        stp_yield();
    }
}

void world(void *null)
{
    int n = 5;
    while (n-- > 0)
    {
        printf("world\n");
        stp_yield();
    } }
```

**Mapping DFGs to Single Processors: Multi-Thread Dynamic Schedule**

```
int main(int argc, char **argv)
{
    stp_init();
    stp_create(hello, 0);
    stp_create(world, 0);
    stp_start();
    return 0;
}
```

To compile and execute: `ex1 ../qt/libstp.a ../qt/libqt.a`

```
gcc -c ex1.c -o
```

```
./ex1
```

```
hello
```

```
world
```

```
hello
```

```
world
```

```
hello
```

```
world\nworld\nworld
```

**Mapping DFGs to Single Processors: Multi-Thread Dynamic Schedule**

A multi-threaded version of the SDF scheduler, using the *fft2 actor*

```
void fft2(actorio_t *g) {  
    int a, b;  
    while (1)  
    {  
        while (fifo_size(g->in[0]) >= 2)  
        {  
            a = get_fifo(g->in[0]);  
            b = get_fifo(g->in[0]);  
            put_fifo(g->out[0], a+b);  
            put_fifo(g->out[0], a-b);  
        }  
        stp_yield();  
    }  
}
```

**Mapping DFGs to Single Processors: Multi-Thread Dynamic Schedule**

```
void main()
{
    fifo_t q1, q2, q3, q4;
    actorio_t fft2_io = {{&q2}, {&q3}};
    ...
    stp_create(fft2, &fft2_io); // create thread
    ...
    stp_start();                // start system scheduler
}
```


Note, as before, the *actor* code must enable convergence to the PASS *firing rate* (through while loops) in order to avoid *queue* overflow

## Mapping DFGs to Single Processors: Static Schedule

From the PASS analysis of an SDF graph, we know at least one solution for a feasible sequential schedule

This solution can be used to optimize the implementation in several ways

- We can remove the *firing rules* since we know the exact sequential schedule This yields only a small performance benefit
- We can also determine an optimal interleaving of the *actors* to minimize the storage requirements for the *queues*
- Finally, we can create a fully **inlined** version of the SDF graph which eliminates the *queues* altogether



```
while(1) {  
    // call A four times  
    A(); A(); A(); A();  
  
    // call B two times  
    B(); B();  
  
    // call C one time  
    C();  
}
```

Here, the relative *firing rates* of A, B, and C must be 4, 2, and 1 to yield a PASS

**Mapping DFGs to Single Processors: Static Schedule**

Given the interleaving schedule on the right, *queue* AB will need to store at most 4 *tokens* and *queue* BC at most 2 *tokens* in steady-state

However, the interleaving schedule (A,A,B,A,A,B,C) is better because the maximum # of tokens on any *queue* is now 2

Therefore, the schedule determined using PASS is **not** necessarily the optimal (in fact, finding the best schedule is an optimization problem)

As noted, implementing a truly static schedule means we do NOT need to check *firing rules* since the required tokens are guaranteed to be present

Consider optimizing the four-point FFT with a **single-thread** SDF system and a **static schedule**

The 3 *actors*, *reorder*, *fft2* and *fft4mag*, have firing rates 1, 2 and 1, which yields a static, cyclic schedule [*reorder*, *fft2*, *fft2*, *fft4mag*]



**Software Implementation: Sequential Targets with Static Schedule**

There are two simple **optimizations** that can be applied here

- The *firing schedule* is **static** and **fixed**, and therefore the access order of *queues* is also fixed

This allows the queues to be *optimized out* and replaced with **fixed variables**

The *queue* access can be replaced as shown in the comments

```
loop {  
    ...  
    q1.put(value1); // replace with r1 = value1;  
    q1.put(value2); // replace with r2 = value2;  
    ...  
    .. = q1.get(); // replace with .. = r1;  
    .. = q1.get(); // replace with .. = r2;  
}
```

**Software Implementation: Sequential Targets with Static Schedule**

- A second optimization involves **inline**'ing the *actor* code in the main program  
In combination with the above optimization, this *eliminates* the *firing rules* and reduces the entire dataflow graph to a **single** function

```
void dftsystem(int in0, in1, in2, in3,  
               *out0, *out1, *out2, *out3) {  
    int reorder_out0, reorder_out1;  
    int reorder_out2, reorder_out3;  
    int fft2_0_out0, fft2_0_out1;  
    int fft2_0_out2, fft2_0_out3;  
    int fft2_1_out0, fft2_1_out1;  
    int fft2_1_out2, fft2_1_out3;  
    int fft4mag_0_out0, fft4mag_0_out1;  
    int fft4mag_0_out2, fft4mag_0_out3;  
  
    // Reorder operation  
    reorder_out0 = in0; reorder_out1 = in2;  
    reorder_out2 = in1; reorder_out3 = in3;
```

**Software Implementation: Sequential Targets with Static Schedule**

```
// Two fft2 implementations
fft2_0_out0 = reorder_out0 + reorder_out1;
fft2_0_out1 = reorder_out0 - reorder_out1;
fft2_1_out0 = reorder_out2 + reorder_out3;
fft2_1_out1 = reorder_out2 - reorder_out3;

// fft4 implementation
fft4mag_out0 = (fft2_0_out0 + fft2_1_out0) *
               (fft2_0_out0 + fft2_1_out0);
fft4mag_out1 = (fft2_0_out1 * fft2_0_out1) -
               (fft2_1_out1 * fft2_1_out1);
fft4mag_out2 = (fft2_0_out0 - fft2_1_out0) *
               (fft2_0_out0 - fft2_1_out0);
fft4mag_out3 = (fft2_0_out1 * fft2_0_out1) -
               (fft2_1_out1 * fft2_1_out1);
```

**Software Implementation: Sequential Targets with Static Schedule**

These optimizations reduce the runtime of the program significantly

For example, we have eliminated testing of the *firing rules* and calls to the *queue* and *actor* functions

This is possible here because a *valid PASS* could be determined from the DFG, as well as *fixed schedule* to implement the PASS

Note that we have traded some of the **runtime flexibility** for **improved efficiency**